

```

import numpy as np
import pandas as pd
from sklearn import *
import matplotlib.pyplot as plt

data = datasets.load_diabetes()

x = pd.DataFrame(data['data'], columns=data['feature_names'])
x['patient_id'] = np.arange(x.shape[0])
x['species'] = 1

y = data['target']
x['y'] = y # Oops...

```

x

	age	sex	bmi	bp	s1	s2	s3
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142
..
437	0.041708	0.050680	0.019662	0.059744	-0.005697	-0.002566	-0.028674
438	-0.005515	0.050680	-0.015906	-0.067642	0.049341	0.079165	-0.028674
439	0.041708	0.050680	-0.015906	0.017293	-0.037344	-0.013840	-0.024993
440	-0.045472	-0.044642	0.039062	0.001215	0.016318	0.015283	-0.028674
441	-0.045472	-0.044642	-0.073030	-0.081413	0.083740	0.027809	0.173816

	s4	s5	s6	patient_id	species	y
0	-0.002592	0.019907	-0.017646	0	1	151.0
1	-0.039493	-0.068332	-0.092204	1	1	75.0
2	-0.002592	0.002861	-0.025930	2	1	141.0
3	0.034309	0.022688	-0.009362	3	1	206.0
4	-0.002592	-0.031988	-0.046641	4	1	135.0
..
437	-0.002592	0.031193	0.007207	437	1	178.0
438	0.034309	-0.018114	0.044485	438	1	104.0

```
439 -0.011080 -0.046883 0.015491      439      1 132.0
440  0.026560  0.044529 -0.025930      440      1 220.0
441 -0.039493 -0.004222  0.003064      441      1  57.0

[442 rows x 13 columns]

x_train, x_test, y_train, y_test = model_selection.train_test_split(x,
y, random_state=42)

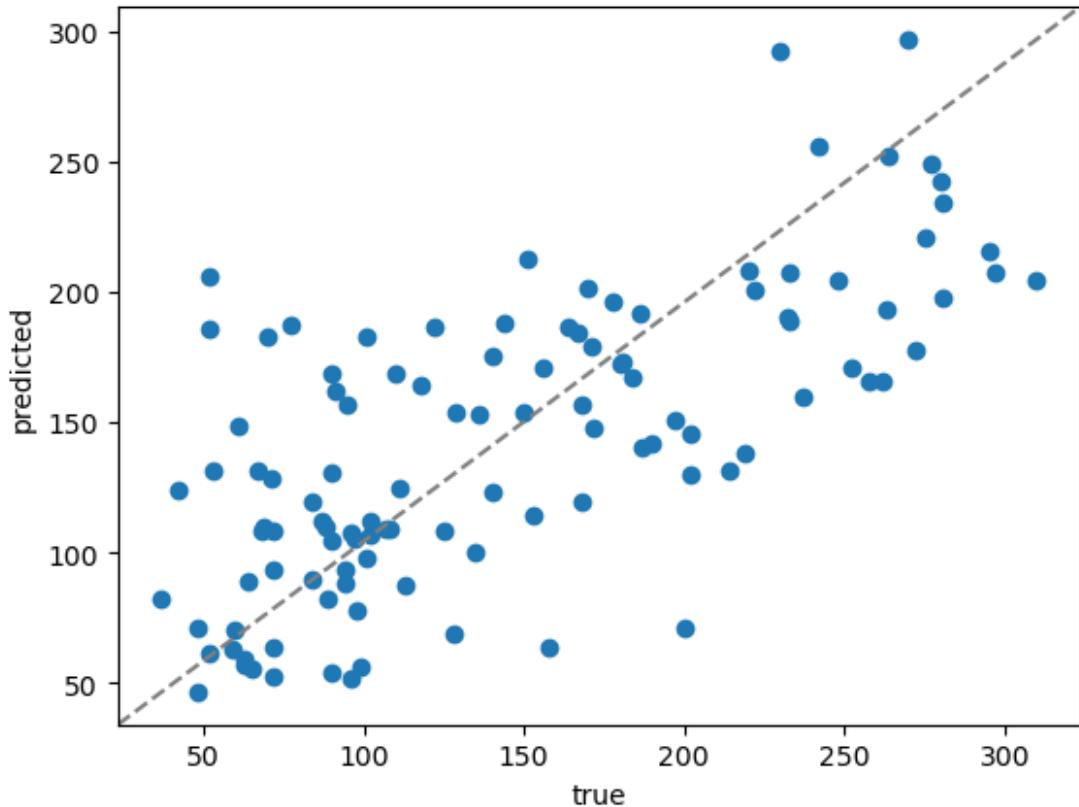
# columns = x.columns
columns = x.columns.difference({'y', 'patient_id', 'species'})
x_train = x_train[columns]
x_test = x_test[columns]

# regressor = svm.SVR(**{'C': 700, 'kernel': 'linear'})
regressor = linear_model.LinearRegression()
regressor.fit(x_train, y_train)
y_hat = regressor.predict(x_test)
print(regressor.score(x_test, y_test))

plt.scatter(y_test, y_hat)
plt.xlabel('true')
plt.ylabel('predicted')
xlim = plt.xlim()
ylim = plt.ylim()
plt.plot(xlim, ylim, '--', c='grey')
plt.xlim(xlim)
plt.ylim(ylim)

0.4849058889476757

(34.09984366041451, 309.2288595404606)
```



```
x_filtered =
feature_selection.VarianceThreshold(threshold=0.2).fit_transform(x_train)
regressor.fit(x_filtered, y_train)
y_hat = regressor.predict(x_test)
print(regressor.score(x_test, y_test))

plt.scatter(y_test, y_hat)
plt.xlabel('true')
plt.ylabel('predicted')
xlim = plt.xlim()
ylim = plt.ylim()
plt.plot(xlim, ylim, '--', c='grey')
plt.xlim(xlim)
plt.ylim(ylim)

-----
-----
ValueError                                     Traceback (most recent call
last)
Cell In [28], line 1
----> 1 x_filtered =
feature_selection.VarianceThreshold(threshold=0.2).fit_transform(x_train)
```

```

2 regressor.fit(x_filtered, y_train)
3 y_hat = regressor.predict(x_test)

File /usr/lib/python3/dist-packages/sklearn/utils/_set_output.py:142,
in _wrap_method_output.<locals>.wrapped(self, X, *args, **kwargs)
  140 @wraps(f)
  141 def wrapped(self, X, *args, **kwargs):
--> 142     data_to_wrap = f(self, X, *args, **kwargs)
  143     if isinstance(data_to_wrap, tuple):
  144         # only wrap the first output for cross decomposition
  145         return (
  146             _wrap_data_with_container(method, data_to_wrap[0],
X, self),
  147             *data_to_wrap[1:],
  148         )
  149     else:
  150         return _wrap_data_with_container(method, data_to_wrap)

File /usr/lib/python3/dist-packages/sklearn/base.py:859, in
TransformerMixin.fit_transform(self, X, y, **fit_params)
  855 # non-optimized default implementation; override when a better
  856 # method is possible for a given clustering algorithm
  857 if y is None:
  858     # fit method of arity 1 (unsupervised transformation)
--> 859     return self.fit(X, **fit_params).transform(X)
  860 else:
  861     # fit method of arity 2 (supervised transformation)
  862     return self.fit(X, y, **fit_params).transform(X)

File
/usr/lib/python3/dist-packages/sklearn/feature_selection/_variance_thr
eshold.py:125, in VarianceThreshold.fit(self, X, y)
  123     if X.shape[0] == 1:
  124         msg += " (X contains only one sample)"
--> 125     raise ValueError(msg.format(self.threshold))
  126 return self

ValueError: No feature in X meets the variance threshold 0.20000

gcv = model_selection.GridSearchCV(
    regressor, {
        # 'n_estimators': np.linspace(10, 100, 9, endpoint=True,
        dtype=int),
        'criterion': ['squared_error', 'absolute_error',
'friedman_mse', 'poisson'],
        'max_depth': np.arange(1, 5),
        'max_features': ['sqrt', 'log2', None],
    },
    cv=4)
gcv.fit(x_train, y_train)
print(gcv.best_params_)
print(gcv.best_score_)

```

```
{'criterion': 'poisson', 'max_depth': 4, 'max_features': 'sqrt'}  
0.40225974404591913  
  
regressor = svm.SVR()  
gcv = model_selection.GridSearchCV(  
    regressor, [  
        {  
            'kernel': ['linear'],  
            # 'C': [ 1e2, 1e3, 1e5, 1e7],  
            'C': np.arange(100, 2100, 100),  
        },  
        # {  
        #     'kernel': ['poly'],  
        #     'degree': np.arange(1, 5),  
        #     'gamma': ['auto', 'scale'],  
        #     'coef0': [0, 1e-2, 1e-1, 1],  
        #     'C': [1e-2, 1e-1, 1, 1e1, 1e2, 1e3],  
        # },  
        # {  
        #     'kernel': ['sigmoid'],  
        #     'gamma': ['auto', 'scale'],  
        #     'coef0': [0, 1e-2, 1e-1, 1],  
        #     'C': [1e-2, 1e-1, 1, 1e1, 1e2, 1e3],  
        # },  
        # {  
        #     'kernel': ['rbf'],  
        #     'gamma': ['auto', 'scale'],  
        #     'C': [1e-2, 1e-1, 1, 1e1, 1e2, 1e3],  
        # },  
    ],  
    cv=4, n_jobs=8)  
gcv.fit(x_train, y_train)  
print(gcv.best_params_)  
print(gcv.best_score_)  
  
{'C': 700, 'kernel': 'linear'}  
0.4419789143972433
```