

Introduction

Up to this point, we've covered the foundational concepts of machine learning. You may have already experimented with various hyperparameters to optimize model performance. In this notebook, you'll be introduced to several advanced techniques designed to further enhance your models.

This notebook consists the following parts:

- A: Data retrieval
 - B: Feature engineering
 - C: Model evaluation
 - D: Ensemble learning
 - E: Pipelines
 - F: Bring it all together
 - G: Bonus: ML OPS
-

Learning Objectives

By the end of this two weeks you will be able to:

- Understand the fundamental concepts of ensemble learning.
 - Use evaluation techniques to assess models performance.
 - Enhance model performance by feature engineering.
 - Build pipelines for model development and preprocessing
-

Instructions

- Ensure you fully understand the requirements and objectives of the assignment.
- Review the notebooks referred in the tasks
- If you need additional context or clarification, please check the provided videos or background literature.
- Work through each part of the assignment methodically, ensuring all tasks are completed.
- Update your repository with your new created work

Additional Notes:

- Do not add datafiles to your repository. Repositories with datafiles will not be accepted
- Class solutions should be delivered in python files. Not in notebooks
- When AI tools are used, you must provide proper references and explanations for how they were utilized. Failure to do so will be considered as academic fraud
- The bonus assignment are not mandatory

- Use PEP8

Good luck!

F.Feenstra

Part A. Data retrieval

The dataset you can use for this notebook is the lung dataset from Maastricht University. It comprises 89 non-small cell lung cancer (NSCLC) patients records who underwent surgical treatment. The study where the data is from explored the relationship between radiomic imaging features and gene expression profiles. The samples were collected through biopsies at MAASTRO Clinic in The Netherlands, and the dataset is publicly available.

The authors of the related paper discovered that a prognostic radiomic signature, which captures intra-tumor heterogeneity, is closely associated with underlying gene expression patterns. Developing a machine learning model to predict histology from the Clinical and Genetic Lung data can improve diagnostic accuracy and treatment personalization. In this notebook we will develop such a prediction model.

Available Datasets:

- Lung metadata dataset [1]
- Gene expression dataset [2]

Important It is also allowed to use your own dataset from your own project if this data is highly dimensional and contains genetic information.

[1] NSCLC-Radiomics-Genomics

[2] Gene Expression Data - GSE58661

[3] Aerts HJWL, Rios Velazquez E, Leijenaar RTH, Parmar C, Grossmann P, Carvalho S, Bussink J, Monshouwer R, Haibe-Kains B, Rietveld D, Hoebers F, Rietbergen MM, Leemans CR, Dekker A, Quackenbush J, Gillies RJ, & Lambin P. (2015). Data From NSCLC-Radiomics-Genomics. The Cancer Imaging Archive. <https://doi.org/10.7937/K9/TCIA.2015.L4FRET6Z>

Data retrieval task

- Retrieve the data. (No cleaning needed yet)

```
import pandas as pd
import gzip

lung_pd = pd.read_excel('Lung3.metadata.xls')
lung_pd.rename(columns=lambda x: x.replace('characteristics.tag.', ''),
               inplace=True)
print(lung_pd.columns)
cols = ['title', 'histology', 'tumor.size.maximumdiameter',
```

```

'stage.primary.tumor', 'gender']
lung_pd = lung_pd[cols]
lung_pd.dropna(inplace=True)
print(lung_pd.columns)
print(lung_pd.head())

lung_pd['histology'] = lung_pd['histology'].apply(lambda x:
x.replace('\xa0', ' ').lower().strip())
uniques = set(list(lung_pd['histology']))

data_gse = pd.read_csv('GSE58661_series_matrix.txt.gz', sep='\t',
comment='!', quotechar='', compression='gzip')
data_gse.dropna(inplace=True)
data_gse.set_index("ID_REF", inplace=True)
print(data_gse.columns)
data_gse.head()

Index(['sample.name', 'title', 'CEL.file', 'source.location',
'organism',
       'gender', 'histology', 'tumor.size.maximumdiameter',
       'stage.primary.tumor', 'stage.nodes', 'stage.mets',
'primaryVSmets',
       'grade', 'molecule tested', 'label', 'platform'],
      dtype='object')
Index(['title', 'histology', 'tumor.size.maximumdiameter',
       'stage.primary.tumor', 'gender'],
      dtype='object')
   title           histology  tumor.size.maximumdiameter
\
0  lung_1  Squamous Cell Carcinoma, NOS          4.0
1  lung_2  Adenocarcinoma, Papillary, NOS         1.3
2  lung_3                  Non-Small Cell        11.0
4  lung_5  Squamous Cell Carcinoma, NOS          7.8
5  lung_6  Adenocarcinoma, NOS                     3.5

  stage.primary.tumor gender
0                  pT2     M
1                  pT1     M
2                  pT3     M
4                  pT3     F
5                  pT2     M
Index(['GSM1416528', 'GSM1416529', 'GSM1416530', 'GSM1416531',
'GSM1416532',
       'GSM1416533', 'GSM1416534', 'GSM1416535', 'GSM1416536',

```

```

'GSM1416537',
    'GSM1416538', 'GSM1416539', 'GSM1416540', 'GSM1416541',
'GSM1416542',
    'GSM1416543', 'GSM1416544', 'GSM1416545', 'GSM1416546',
'GSM1416547',
    'GSM1416548', 'GSM1416549', 'GSM1416550', 'GSM1416551',
'GSM1416552',
    'GSM1416553', 'GSM1416554', 'GSM1416555', 'GSM1416556',
'GSM1416557',
    'GSM1416558', 'GSM1416559', 'GSM1416560', 'GSM1416561',
'GSM1416562',
    'GSM1416563', 'GSM1416564', 'GSM1416565', 'GSM1416566',
'GSM1416567',
    'GSM1416568', 'GSM1416569', 'GSM1416570', 'GSM1416571',
'GSM1416572',
    'GSM1416573', 'GSM1416574', 'GSM1416575', 'GSM1416576',
'GSM1416577',
    'GSM1416578', 'GSM1416579', 'GSM1416580', 'GSM1416581',
'GSM1416582',
    'GSM1416583', 'GSM1416584', 'GSM1416585', 'GSM1416586',
'GSM1416587',
    'GSM1416588', 'GSM1416589', 'GSM1416590', 'GSM1416591',
'GSM1416592',
    'GSM1416593', 'GSM1416594', 'GSM1416595', 'GSM1416596',
'GSM1416597',
    'GSM1416598', 'GSM1416599', 'GSM1416600', 'GSM1416601',
'GSM1416602',
    'GSM1416603', 'GSM1416604', 'GSM1416605', 'GSM1416606',
'GSM1416607',
    'GSM1416608', 'GSM1416609', 'GSM1416610', 'GSM1416611',
'GSM1416612',
    'GSM1416613', 'GSM1416614', 'GSM1416615', 'GSM1416616'],
dtype='object')

```

	GSM1416528	GSM1416529	GSM1416530	GSM1416531
GSM1416532 \ ID_REF				
AFFX-BioB-3_at 7.211031	7.376915	8.024915	7.522543	7.152864
AFFX-BioB-5_at 6.753131	6.984530	7.427048	7.077207	6.849513
AFFX-BioB-M_at 7.077163	7.330576	8.010530	7.334551	7.143286
AFFX-BioC-3_at 10.701328	10.922741	11.390638	10.936703	10.791909
AFFX-BioC-5_at 10.823792	11.032030	11.533338	11.018510	10.875259

GSM1416533	GSM1416534	GSM1416535	GSM1416536
------------	------------	------------	------------

GSM1416537	\
ID_REF	
AFFX-BioB-3_at	7.704962
8.141578	7.779374
AFFX-BioB-5_at	7.195947
7.632631	7.429800
AFFX-BioB-M_at	7.753484
7.992361	7.712429
AFFX-BioC-3_at	11.221732
11.412668	11.084229
AFFX-BioC-5_at	11.276973
11.470336	10.588372
	10.628909
	10.700420
	10.779332
	... GSM1416607
	GSM1416608
	GSM1416609
	GSM1416610 \
ID_REF	
AFFX-BioB-3_at	...
	7.611992
AFFX-BioB-5_at	...
	7.175015
AFFX-BioB-M_at	...
	7.612393
AFFX-BioC-3_at	...
	11.097486
AFFX-BioC-5_at	...
	11.265739
	7.345206
	6.913879
	7.324635
	10.908682
	10.781228
	10.935723
	10.900765
	11.101546
	GSM1416611
	GSM1416612
	GSM1416613
	GSM1416614
GSM1416615	\
ID_REF	
AFFX-BioB-3_at	7.756286
7.039592	7.215856
AFFX-BioB-5_at	7.347570
6.575376	6.740992
AFFX-BioB-M_at	7.745624
6.975162	7.152722
AFFX-BioC-3_at	11.200635
10.367009	10.707980
AFFX-BioC-5_at	11.343329
10.494806	10.736531
	10.770440
	10.900509
	10.838008
	GSM1416616
ID_REF	
AFFX-BioB-3_at	7.393667
AFFX-BioB-5_at	7.041106
AFFX-BioB-M_at	7.366620
AFFX-BioC-3_at	10.753221
AFFX-BioC-5_at	10.863600

[5 rows x 89 columns]

Part B. Feature engineering

Mind you that choosing an algorithm and hyperparameter tuning might not be enough. If your data is of low quality, the algorithm will have poor performance as well. This is where feature engineering comes into play. Feature engineering involves transforming raw data into meaningful features that better represent the underlying problem to the predictive models, ultimately enhancing the model's performance. By selecting, creating, and refining features, you ensure that your data highlights the most relevant patterns and relationships, allowing the algorithm to learn more effectively. Proper feature engineering can often make the difference between a moderate model and a highly accurate one, even more so than the choice of algorithm itself. Possible modifications:

- creation of new features derived from original features
- selection of features
- encoding features
- log transformation
- scaling
- dimension reduction (*e.g.* PCA)

Feature engineering Task

- Review the [Study Case Feature Engineering notebook](#).
- Review the [Study Case RNA-seq Preparation notebook](#).
- Assess which data preparation and feature engineering steps could be beneficial for your dataset.
- Implement a `DataProcessor` class tailored to your dataset and test it in the cell below. **Make sure it's an appropriate `sklearn` object, such as a `Transformer` or `Pipeline`.**
- Update your repository with a new directory named `optimization`, including the following:
 - The `DataProcessor` class as a Python file, complete with thorough documentation.
 - An evaluation document that details and justifies your choices using a well-reasoned, argumentative approach.

```
# YOUR CODE HERE TO DEMONSTRATE THE USAGE OF THE PREPROCESSOR CLASS
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import numpy as np

data_gse_trans = data_gse.transpose()

# To combat skewed data
data_gse_log = np.log1p(data_gse_trans)

# Standardizing since genes can have very different scales of expression.
scaler = StandardScaler()
scaled_gse_trans_log = scaler.fit_transform(data_gse_log)
```

```

# Performing PCA I know there is a better way to do like doing a pca
# on all data and then look at what point it becomes more than 0.5
# total.
# But this is to practice recursion.
def find_components(data, start=2):
    pca = PCA(n_components=start)
    pca_res = pca.fit_transform(data)
    expl_var_list = pca.explained_variance_ratio_

    if sum(expl_var_list) >= 0.5:
        print(f"N component to reach 50% explained of the variance:
{start}")
        print(f"Explained variance list: {expl_var_list}")
        return start, pca, expl_var_list, pca_res
    else:
        start += 1
        return find_components(data, start)

start, pca, expl_var_list, pca_res =
find_components(scaled_gse_trans_log)
columns1 = [f"PCA_{i+1}" for i in range(start)]

pca_res_df = pd.DataFrame(data=pca_res, columns=columns1)

merged_df = pd.merge(lung_pd, pca_res_df, left_index=True,
right_index=True)
print(merged_df.head())

N component to reach 50% explained of the variance: 16
Explained variance list: [0.09358163 0.08701857 0.04763506 0.0439992
0.03330682 0.02687025
0.02465527 0.02241783 0.01850185 0.01725881 0.01621868 0.01531552
0.01433772 0.01354639 0.01325415 0.01315874]
      title           histology  tumor.size.maximumdiameter
\\
0  lung_1  squamous cell carcinoma, nos          4.0
1  lung_2  adenocarcinoma, papillary, nos         1.3
2  lung_3                  non-small cell         11.0
4  lung_5  squamous cell carcinoma, nos          7.8
5  lung_6  adenocarcinoma, nos                   3.5

stage.primary.tumor gender      PCA_1      PCA_2      PCA_3
PCA_4 \
0             pT2      M -12.563698  143.350916 -74.113707 -

```

```

29.447517
1          pT1      M -14.585510 -136.108180 -31.682665 -
93.293509
2          pT3      M -72.359515   31.080436 -7.212311 -
23.003114
4          pT3      F -61.898539 -2.743114 -34.230067
42.663674
5          pT2      M  17.562691 -86.546114 -10.003988
41.723127

    PCA_5 ... PCA_7 PCA_8 PCA_9 PCA_10
PCA_11 \
0 -22.873815 ... 32.850302 -33.709224 62.696240 -4.411993
19.968412
1 -35.638663 ... 13.783358  6.443782 -3.767860 52.725018
55.846064
2 -3.620890 ... -45.965215 -77.529470 83.729802 -1.545160
119.487783
4 44.716229 ... -50.697751  1.434348 -16.768901 -27.776140 -
22.891227
5 17.580223 ... -24.641162  0.501947 26.509949 -18.814753
5.222445

    PCA_12 PCA_13 PCA_14 PCA_15 PCA_16
0 -11.276013  6.471942 -0.184468 -22.009904  5.381011
1 -2.812119 -3.451776  81.455464 12.315144 -109.935105
2  7.006954 -42.027303 -7.773518 -64.385517  54.684574
4 -16.557359 -26.567846 29.651508 -16.065088 -10.076500
5 21.309420  2.921614  8.611129 -11.747633 -9.636650

```

[5 rows x 21 columns]

```

# Encoder used for turning strings/classifications into numeric values
from sklearn.preprocessing import LabelEncoder

def encode_column(df, column_enc):
    encoder = LabelEncoder()
    df[column_enc] = encoder.fit_transform(df[column_enc])
    return df

merged_df = encode_column(merged_df, 'histology')
merged_df = encode_column(merged_df, 'stage.primary.tumor')

from DataProcessor import DataProcessor

pipeline = DataProcessor(input_file_exp =
    'GSE58661_series_matrix.txt.gz', input_file_meta =
    "Lung3.metadata.xls",
                           cols_to_grab = ['title',
    'characteristics.tag.gender', 'characteristics.tag.histology'],

```

```

        'characteristics.tag.tumor.size.maximumdiameter'],
target='characteristics.tag.histology')

pipeline_finished, encoders = pipeline.pipeline_run()
print(encoders)
print(pipeline_finished.head())
print(pipeline.look_at_encode(encoders, "title"))

N component to reach 50% explained of the variance: 14
Explained variance list: [0.13625197 0.0657839  0.04725548 0.04508164
0.03637912 0.02774053
0.02481753 0.02132643 0.02092515 0.01706445 0.01596801 0.01559264
0.01437887 0.01392811]
['title', 'gender', 'histology']
{'title': LabelEncoder(), 'gender': LabelEncoder()}
    title   gender           histology
tumor.size.maximumdiameter \
0      0       1  Squamous Cell Carcinoma, NOS
4.0
1      11      1  Adenocarcinoma, Papillary, NOS
1.3
2      22      1           Non-Small Cell
11.0
4      43      0  Squamous Cell Carcinoma, NOS
7.8
5      54      1  Adenocarcinoma, NOS
3.5

PCA_1     PCA_2     PCA_3     PCA_4     PCA_5     PCA_6
PCA_7 \
0  15.075208 -5.472546  0.145589   2.995557  3.196469 -0.816480 -
8.442277
1  -6.797299  9.107596 -4.149766  10.091782 -3.867643  2.197285 -
0.667493
2   7.656499  5.641671  2.394907   0.509648 -0.128076  7.513621 -
9.243774
4   6.930568 -1.981705 -3.092332  -2.864091 -4.972384  0.097903
3.851141
5  -6.866977  0.155878 -4.951343  -4.213497  0.500106 -1.366388
0.556406

PCA_8     PCA_9     PCA_10    PCA_11    PCA_12    PCA_13
PCA_14
0  2.144322  0.371191 -1.495906 -4.127096 -0.762431 -3.680627 -
1.920025
1  2.812155 -2.059079 -2.073999  5.962730  0.650969  3.316948 -
5.870156
2  5.257675 -4.131223  5.760677 -8.454589  1.021437 12.263974

```

```

9.003247
4 -4.317892 4.827766 1.086183 1.126033 3.397036 1.218709
3.900536
5 -2.570907 -0.207244 0.158308 -1.783719 -2.991035 5.627051 -
0.053776
{'lung_1': 0, 'lung_10': 1, 'lung_11': 2, 'lung_12': 3, 'lung_13': 4,
'lung_14': 5, 'lung_15': 6, 'lung_16': 7, 'lung_17': 8, 'lung_18': 9,
'lung_19': 10, 'lung_2': 11, 'lung_20': 12, 'lung_21': 13, 'lung_22': 14,
'lung_23': 15, 'lung_24': 16, 'lung_25': 17, 'lung_26': 18,
'lung_27': 19, 'lung_28': 20, 'lung_29': 21, 'lung_3': 22, 'lung_30': 23,
'lung_31': 24, 'lung_32': 25, 'lung_33': 26, 'lung_34': 27,
'lung_35': 28, 'lung_36': 29, 'lung_37': 30, 'lung_38': 31, 'lung_39': 32,
'lung_40': 33, 'lung_41': 34, 'lung_42': 35, 'lung_43': 36,
'lung_44': 37, 'lung_45': 38, 'lung_46': 39, 'lung_47': 40, 'lung_48': 41,
'lung_49': 42, 'lung_5': 43, 'lung_50': 44, 'lung_51': 45,
'lung_52': 46, 'lung_53': 47, 'lung_54': 48, 'lung_55': 49, 'lung_56': 50,
'lung_57': 51, 'lung_58': 52, 'lung_59': 53, 'lung_6': 54,
'lung_60': 55, 'lung_61': 56, 'lung_62': 57, 'lung_63': 58, 'lung_64': 59,
'lung_65': 60, 'lung_66': 61, 'lung_67': 62, 'lung_68': 63,
'lung_69': 64, 'lung_7': 65, 'lung_70': 66, 'lung_71': 67, 'lung_72': 68,
'lung_73': 69, 'lung_74': 70, 'lung_75': 71, 'lung_76': 72,
'lung_77': 73, 'lung_78': 74, 'lung_79': 75, 'lung_8': 76, 'lung_80': 77,
'lung_81': 78, 'lung_82': 79, 'lung_83': 80, 'lung_84': 81,
'lung_85': 82, 'lung_86': 83, 'lung_87': 84, 'lung_88': 85, 'lung_89': 86,
'lung_9': 87}
{'lung_1': 0, 'lung_10': 1, 'lung_11': 2, 'lung_12': 3, 'lung_13': 4,
'lung_14': 5, 'lung_15': 6, 'lung_16': 7, 'lung_17': 8, 'lung_18': 9,
'lung_19': 10, 'lung_2': 11, 'lung_20': 12, 'lung_21': 13, 'lung_22': 14,
'lung_23': 15, 'lung_24': 16, 'lung_25': 17, 'lung_26': 18,
'lung_27': 19, 'lung_28': 20, 'lung_29': 21, 'lung_3': 22, 'lung_30': 23,
'lung_31': 24, 'lung_32': 25, 'lung_33': 26, 'lung_34': 27,
'lung_35': 28, 'lung_36': 29, 'lung_37': 30, 'lung_38': 31, 'lung_39': 32,
'lung_40': 33, 'lung_41': 34, 'lung_42': 35, 'lung_43': 36,
'lung_44': 37, 'lung_45': 38, 'lung_46': 39, 'lung_47': 40, 'lung_48': 41,
'lung_49': 42, 'lung_5': 43, 'lung_50': 44, 'lung_51': 45,
'lung_52': 46, 'lung_53': 47, 'lung_54': 48, 'lung_55': 49, 'lung_56': 50,
'lung_57': 51, 'lung_58': 52, 'lung_59': 53, 'lung_6': 54,
'lung_60': 55, 'lung_61': 56, 'lung_62': 57, 'lung_63': 58, 'lung_64': 59,
'lung_65': 60, 'lung_66': 61, 'lung_67': 62, 'lung_68': 63,
'lung_69': 64, 'lung_7': 65, 'lung_70': 66, 'lung_71': 67, 'lung_72': 68,
'lung_73': 69, 'lung_74': 70, 'lung_75': 71, 'lung_76': 72,
'lung_77': 73, 'lung_78': 74, 'lung_79': 75, 'lung_8': 76, 'lung_80': 77,
'lung_81': 78, 'lung_82': 79, 'lung_83': 80, 'lung_84': 81,
'lung_85': 82, 'lung_86': 83, 'lung_87': 84, 'lung_88': 85, 'lung_89': 86,
'lung_9': 87}

```

Part C. Model Evaluation

Evaluating the performance of a machine learning model goes beyond just looking at accuracy, as accuracy alone can be misleading, especially in cases where the dataset is imbalanced or where different types of errors have different consequences

- **Detecting Overfitting/Underfitting:** A learning curve can help you understand whether your model is overfitting (performing well on training data but poorly on validation data) or underfitting (performing poorly on both training and validation data)
- **ROC:** The ROC curve shows how well your model distinguishes between classes. It helps in selecting the optimal threshold for classification decisions, particularly when the cost of false positives and false negatives differs significantly. ROC curves are often used to compare models.
- **Confusion matrix:** From the confusion matrix, you can derive other important metrics like precision, recall, F1-score, and specificity, which give a better understanding of how your model is performing across different classes

Evaluation Task

- Review the documentation for `sklearn.metrics` and `learning_curve`.
- Select appropriate metrics for your dataset, including accuracy and indicators of overfitting or underfitting.
- Implement functions to compute and assess these metrics. It is allowed to use libraries.
- Add the evaluation functions as a Python module to your repository.
- Update the evaluation documentation to clearly explain your choices for the evaluation metrics and the rationale behind them.

See also: [Model evaluation video](#)

Part D. Ensemble learning

Ensemble learning is a powerful machine learning technique that combines the predictions of multiple models to improve overall performance, robustness, and accuracy. Rather than relying on a single model, ensemble learning methods aggregate the results of several models—often called "weak learners"—to produce a stronger predictive model. The key idea behind ensemble learning is that by combining models, the weaknesses of individual models can be offset, leading to better generalization on unseen data. Popular ensemble ML algorithms are the `Random Forest` and `XGBoost`. Here's an improved version of the introduction:

voting algorithms

Voting algorithms in ensemble learning combine the predictions of multiple classifiers to make a final decision, typically based on the consensus or weighted agreement among models. Two primary types of voting are commonly used: hard voting, where the final prediction is

determined by the majority vote, and soft voting, which uses the weighted average of predicted probabilities to determine the outcome. Several sites explain the hard and soft voting algorithm. A clear explanation can be found on <https://www.baeldung.com/cs/hard-vs-soft-voting-classifiers>

Ensemble Task

- Review the [Study Case notebook on ensemble learning](#).
- Try three different algorithms for classification of your data label.
- Implement a hard and soft voting algorithm for model aggregation.
- Compare the performance of the voting algorithm with that of a boosting or bagging algorithm.
- Update your repository with the voting algorithm class

See also: [ensemle learning video](#)

```
# YOUR CODE HERE WITH THE VOTING ALGORITHMS
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score, classification_report,
precision_score, confusion_matrix, roc_curve, roc_auc_score,
recall_score, f1_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from numpy import mean

print(pipeline_finished.columns)
print(pipeline_finished.head())

svccl = SVC(kernel='sigmoid')
knncl = KNeighborsClassifier(n_neighbors=5)
rfcl = RandomForestClassifier(random_state=29)

models = {"svccl": svccl, "knncl":knncl, "rfcl":rfcl}

y = pipeline_finished['histology']
X = pipeline_finished.drop(columns=['histology', 'title', 'gender'])

print(len(X))

rare_class = y.value_counts()[y.value_counts() < 5].index
X = X[~y.isin(rare_class)]
y = y[~y.isin(rare_class)]
print(len(X))

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=20)

svccl.fit(X_train, y_train)
```

```

knncl.fit(X_train, y_train)
rfcl.fit(X_train, y_train)

svccl_pred = svccl.predict(X_test)
knncl_pred = knncl.predict(X_test)
rfcl_pred = rfcl.predict(X_test)

def get_scores(model_name, y_test, model, X, y):

    model.fit(X_train, y_train)
    model_pred = model.predict(X_test)

    print(f"\n\n\nModel name: {model_name}")
    print(f"Precision: {precision_score(y_test, model_pred,
average='weighted', zero_division=0)}")
    print(f"Accuracy: {accuracy_score(y_test, model_pred)}")
    print(f"Recall: {recall_score(y_test, model_pred,
average='weighted', zero_division=0)}")
    print(f"F1 Score: {f1_score(y_test, model_pred,
average='weighted', zero_division=0)}")

    cross_score = cross_val_score(model, X, y, cv=5)
    print(f"Cross_scores:\nCross val mean: {mean(cross_score)}")
    print(f"Cross val list: {cross_score}")

    print(f"ConfusionMatrix:\n Matrix: {confusion_matrix(y_test,
model_pred)}")
    print(f"Classification report:\nReport
{classification_report(y_test, model_pred, zero_division=0)}")
    return model, model_pred

for model in models:
    get_scores(f"{model}", y_test, models[model], X, y)
    print(model)

```

	title	gender	histology	tumor.size.maximumdiameter
PCA_1				
PCA_2	'PCA_2'	'PCA_3'	'PCA_4'	'PCA_5'
PCA_3				'PCA_6'
PCA_4				'PCA_7'
PCA_5				'PCA_8'
PCA_6				
PCA_7				
PCA_8				
PCA_9				
PCA_10				
PCA_11				
PCA_12				
PCA_13				
PCA_14				
				dtype='object')
	title	gender	histology	
	tumor.size.maximumdiameter	\		
0	0	1	Squamous Cell Carcinoma, NOS	
4.0				
1	11	1	Adenocarcinoma, Papillary, NOS	
1.3				

2	22	1	Non-Small Cell					
11.0								
4	43	0	Squamous Cell Carcinoma, NOS					
7.8								
5	54	1	Adenocarcinoma, NOS					
3.5								
	PCA_1	PCA_2	PCA_3	PCA_4	PCA_5	PCA_6		
PCA_7	\							
0	15.075208	-5.472546	0.145589	2.995557	3.196469	-0.816480	-	
	8.442277							
1	-6.797299	9.107596	-4.149766	10.091782	-3.867643	2.197285	-	
	0.667493							
2	7.656499	5.641671	2.394907	0.509648	-0.128076	7.513621	-	
	9.243774							
4	6.930568	-1.981705	-3.092332	-2.864091	-4.972384	0.097903		
	3.851141							
5	-6.866977	0.155878	-4.951343	-4.213497	0.500106	-1.366388		
	0.556406							
	PCA_8	PCA_9	PCA_10	PCA_11	PCA_12	PCA_13		
PCA_14	\							
0	2.144322	0.371191	-1.495906	-4.127096	-0.762431	-3.680627	-	
	1.920025							
1	2.812155	-2.059079	-2.073999	5.962730	0.650969	3.316948	-	
	5.870156							
2	5.257675	-4.131223	5.760677	-8.454589	1.021437	12.263974		
	9.003247							
4	-4.317892	4.827766	1.086183	1.126033	3.397036	1.218709		
	3.900536							
5	-2.570907	-0.207244	0.158308	-1.783719	-2.991035	5.627051	-	
	0.053776							
88								
65								

```

Model name: svccl
Precision: 0.6791666666666667
Accuracy: 0.8
Recall: 0.8
F1 Score: 0.7345454545454546
Cross_scores:
Cross val mean: 0.8153846153846155
Cross val list: [0.84615385 0.76923077 0.84615385 0.84615385
0.76923077]
ConfusionMatrix:
Matrix: [[10  0  0]
 [ 1  0  2]
 [ 1  0  6]]

```

Classification report:

Report	precision	recall	f1-score	
support				
Adenocarcinoma, NOS	0.83	1.00	0.91	10
Non-Small Cell	0.00	0.00	0.00	3
Squamous Cell Carcinoma, NOS	0.75	0.86	0.80	7
accuracy			0.80	20
macro avg	0.53	0.62	0.57	20
weighted avg	0.68	0.80	0.73	20

svccl

Model name: knncl

Precision: 0.6424242424242423

Accuracy: 0.75

Recall: 0.75

F1 Score: 0.6910714285714286

Cross_scores:

Cross val mean: 0.7692307692307694

Cross val list: [0.76923077 0.69230769 0.84615385 0.84615385

0.69230769]

ConfusionMatrix:

Matrix: [[9 0 1]
[1 0 2]
[1 0 6]]

Classification report:

Report	precision	recall	f1-score	
support				
Adenocarcinoma, NOS	0.82	0.90	0.86	10
Non-Small Cell	0.00	0.00	0.00	3
Squamous Cell Carcinoma, NOS	0.67	0.86	0.75	7
accuracy			0.75	20
macro avg	0.49	0.59	0.54	20
weighted avg	0.64	0.75	0.69	20

knncl

Model name: rfcl

Precision: 0.5186813186813187

Accuracy: 0.55

Recall: 0.55

F1 Score: 0.5041176470588236

```

Cross_scores:
Cross val mean: 0.7538461538461538
Cross val list: [0.84615385 0.69230769 0.84615385 0.69230769
0.69230769]
ConfusionMatrix:
Matrix: [[5 0 5]
[1 0 2]
[1 0 6]]
Classification report:
Report                                                 precision    recall   f1-score
support

          Adenocarcinoma, NOS      0.71      0.50      0.59      10
          Non-Small Cell         0.00      0.00      0.00       3
Squamous Cell Carcinoma, NOS      0.46      0.86      0.60       7

          accuracy                   0.55      20
          macro avg                  0.39      0.45      0.40      20
          weighted avg                0.52      0.55      0.50      20

rfcl

from Evaluation import Evaluator
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

svccl = SVC(kernel='sigmoid', probability=True)
knncl = KNeighborsClassifier(n_neighbors=5)
rfcl = RandomForestClassifier(random_state=29)

models = [svccl, knncl, rfcl]

y = pipeline_finished['histology']
X = pipeline_finished.drop(columns=['histology', 'title', 'gender'])

eval_class = Evaluator(models, X, y)

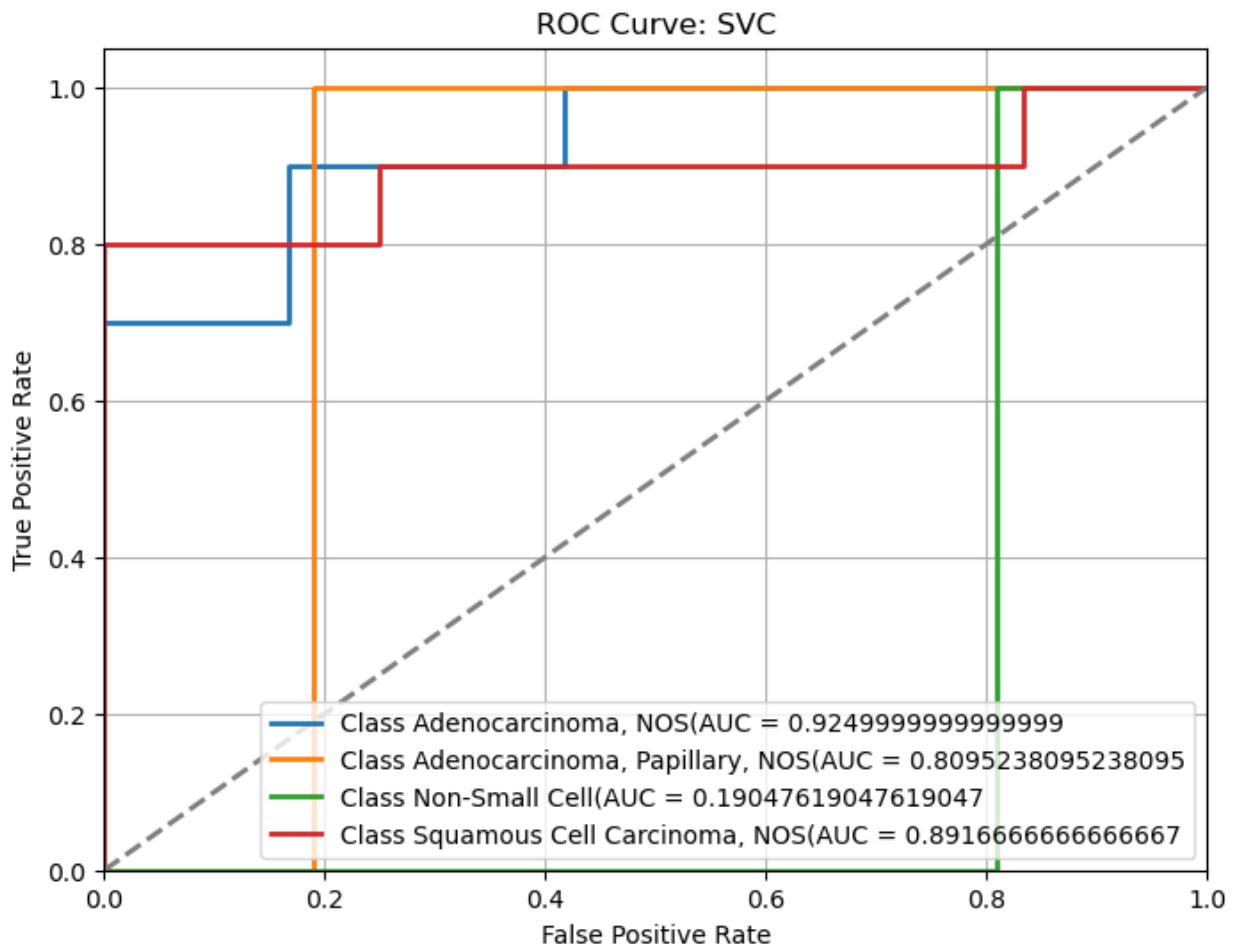
eval_class.execute()

Lost X: 15/88 instances.
Lost X: 15/88 instances.

Model name: SVC
Precision: 0.7045454545454546
Accuracy: 0.7727272727272727
Recall: 0.7727272727272727
F1 Score: 0.7355371900826445

```

```
ConfusionMatrix:  
Matrix: [[9 0 0 1]  
[1 0 0 0]  
[0 0 0 1]  
[2 0 0 8]]  
Classification report:  
Report precision recall f1-score  
support  
  
10 Adenocarcinoma, NOS 0.75 0.90 0.82  
1 Adenocarcinoma, Papillary, NOS 0.00 0.00 0.00  
1 Non-Small Cell 0.00 0.00 0.00  
10 Squamous Cell Carcinoma, NOS 0.80 0.80 0.80  
22 accuracy 0.77  
22 macro avg 0.39 0.43 0.40  
22 weighted avg 0.70 0.77 0.74
```

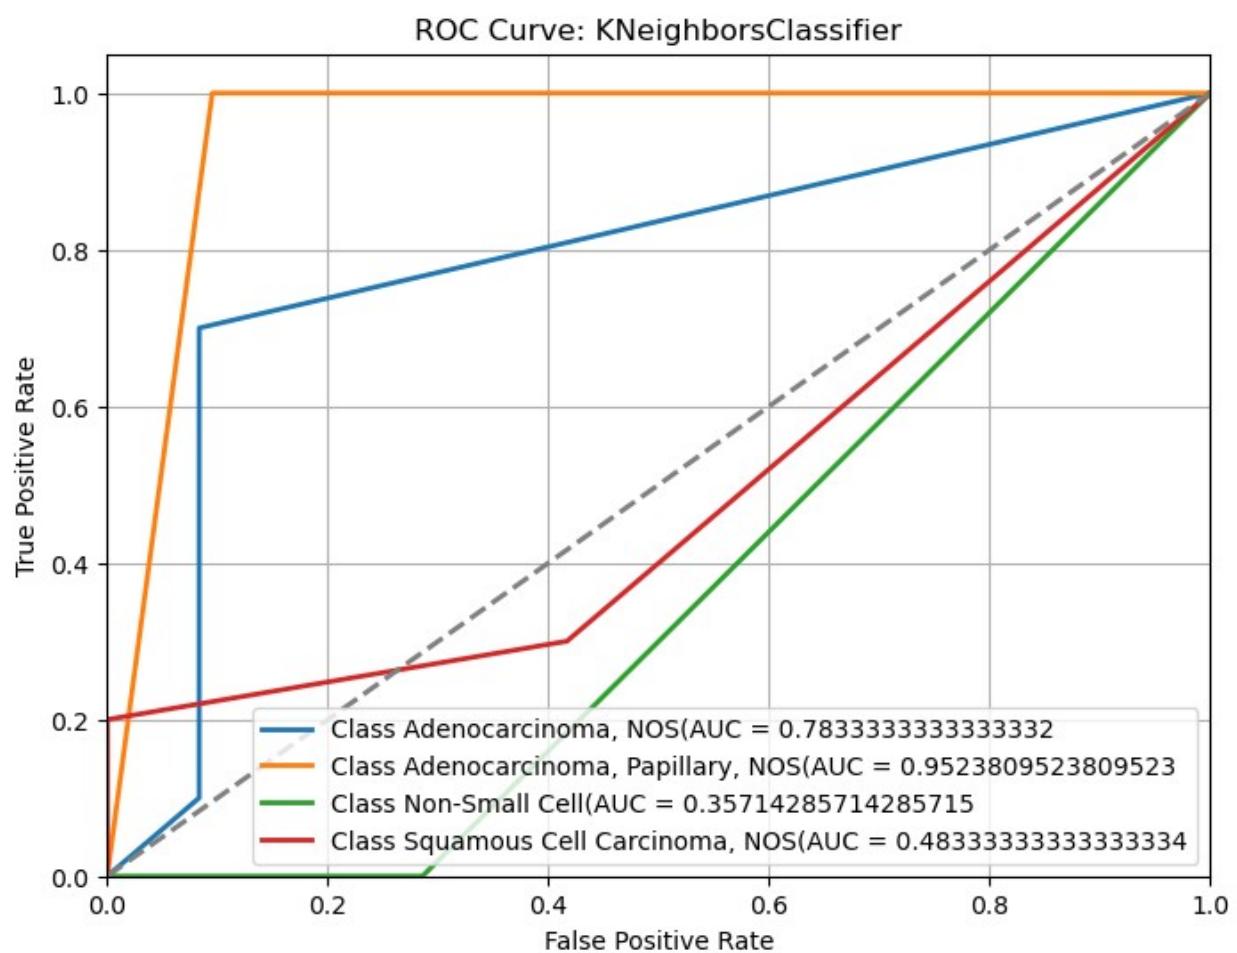


```

Model name: KNeighborsClassifier
Precision: 0.7759412304866852
Accuracy: 0.7727272727272727
Recall: 0.7727272727272727
F1 Score: 0.7723855092276145
ConfusionMatrix:
Matrix: [[0 0 0 0 0]
 [1 8 0 0 1]
 [0 1 0 0 0]
 [0 0 0 0 1]
 [1 0 0 0 9]]
Classification report:
Report                                                 precision
recall   f1-score   support
Adenocarcinoma, Bronchiolo-alveolar Features        0.00      0.00
0.00          0
Adenocarcinoma, NOS                         0.89      0.80

```

0.84	10	Adenocarcinoma, Papillary, NOS	0.00	0.00
0.00	1	Non-Small Cell	0.00	0.00
0.00	1	Squamous Cell Carcinoma, NOS	0.82	0.90
0.86	10			
accuracy				
0.77	22		macro avg	0.34
0.34	22		weighted avg	0.78
0.77	22			0.77



Model name: RandomForestClassifier

Precision: 0.7396694214876033

Accuracy: 0.7727272727272727

Recall: 0.7727272727272727

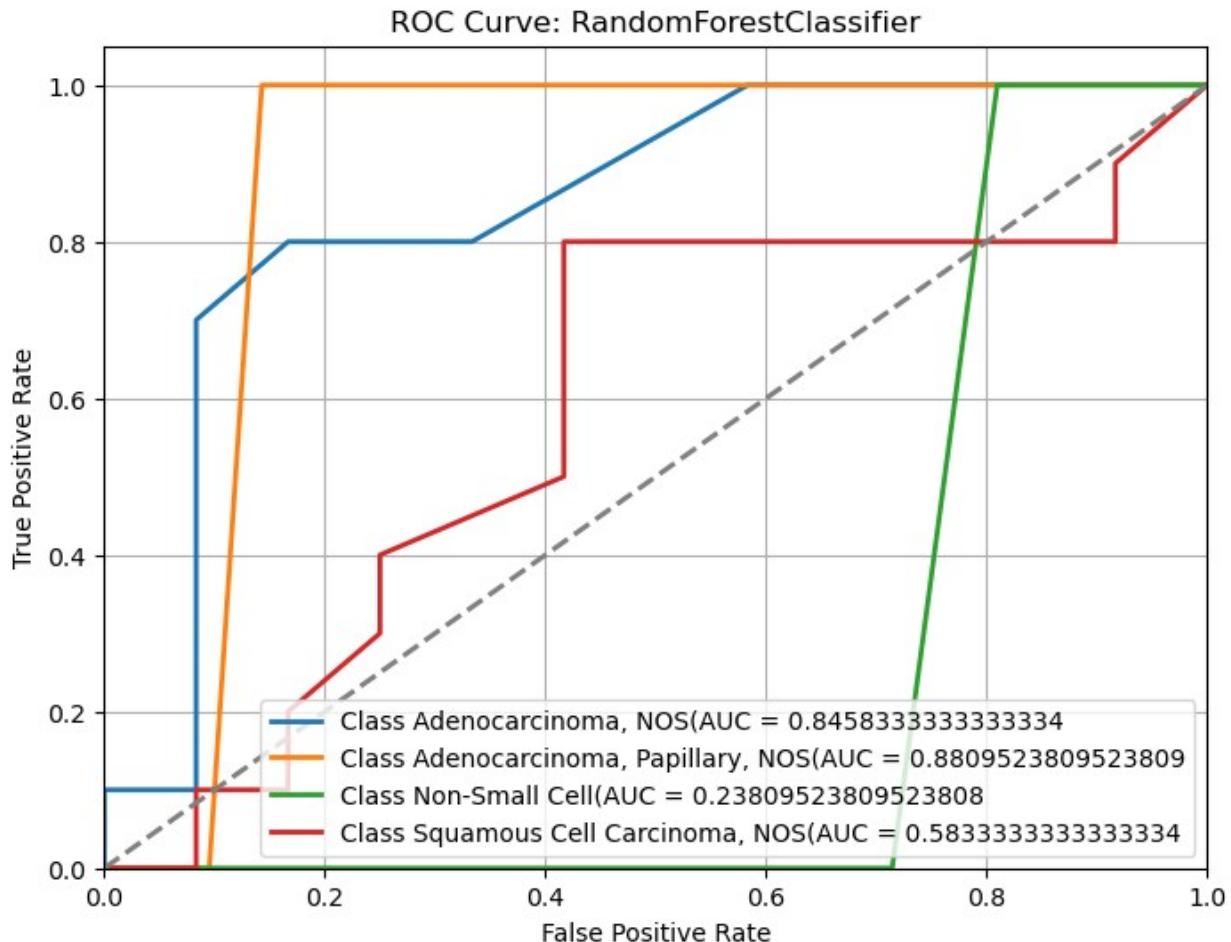
F1 Score: 0.7554112554112554

ConfusionMatrix:

Matrix: [[8 0 0 2]
[1 0 0 0]
[0 0 1 0]
[2 0 0 8]]

Classification report:

		precision	recall	f1-score
	support			
10	Adenocarcinoma, NOS	0.73	0.80	0.76
1	Adenocarcinoma, Papillary, NOS	0.00	0.00	0.00
1	Non-Small Cell	1.00	1.00	1.00
10	Squamous Cell Carcinoma, NOS	0.80	0.80	0.80
22	accuracy			0.77
22	macro avg	0.63	0.65	0.64
22	weighted avg	0.74	0.77	0.76



```
{'SVC': [SVC(kernel='sigmoid', probability=True),
array(['Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
       'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
       'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
       'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
       'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
       'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
       'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
       'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
       'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
       'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS'],
      dtype=object),
<Figure size 800x600 with 1 Axes>],
'KNeighborsClassifier': [KNeighborsClassifier(),
array(['Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
       'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
       'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
       'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS'],
      dtype=object)]}
```

```

        'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
        'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
        'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
        'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
        'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
        'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
        'Adenocarcinoma, Bronchiolo-alveolar Features',
        'Squamous Cell Carcinoma, NOS',
        'Adenocarcinoma, Bronchiolo-alveolar Features',
        'Adenocarcinoma, NOS'], dtype=object),
<Figure size 800x600 with 1 Axes>],
'RandomForestClassifier': [RandomForestClassifier(random_state=29),
array(['Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
        'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
        'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
        'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
        'Non-Small Cell', 'Adenocarcinoma, NOS',
        'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
        'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
        'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
        'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
        'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
        'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS'],
dtype=object),
<Figure size 800x600 with 1 Axes>]}

from scipy.stats import mode

def vote_soft(model_list, X):
    """
    sums the predic_proba across models and pick the highest. models
    must support predict_proba.
    """
    prob_sum = np.sum([model.predict_proba(X) for model in
model_list], axis=0)
    soft_vote_pred = np.argmax(prob_sum, axis=1)
    return soft_vote_pred

def vote_hard(model_list, X):
    #Vote majority across models
    all_preds = np.array([model.predict(X) for model in model_list])
    majority_vote = mode(all_preds, axis=0).mode.flatten()
    return majority_vote

# YOUR CODE HERE TO COMPARE VOTING OUTCOME WITH A BAGGING OR BOOSTING
ALGORITHM FROM SKLEARN

```

```

from VoteClass import Voter
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report

vote = Voter(model_list=models, X_test=eval_class.X_test,
vote_type='s')
results = vote.choose_vote()

boost_model = GradientBoostingClassifier(random_state=42)
boost_model.fit(eval_class.X_train, eval_class.y_train)
boost_pred= boost_model.predict(eval_class.X_test)

print("Boosting Model Results")
print("Accuracy:", accuracy_score(eval_class.y_test, boost_pred))
print("Classification Report:\n",
classification_report(eval_class.y_test, boost_pred, zero_division=0))

print("Vote Model Results")
print("Accuracy:", accuracy_score(eval_class.y_test, results))
print("Classification Report:\n",
classification_report(eval_class.y_test, results, zero_division=0))

Boosting Model Results
Accuracy: 0.5909090909090909
Classification Report:
              precision    recall
f1-score      support
Adenocarcinoma, Bronchiolo-alveolar Features      0.00      0.00
0.00          0
                                         Adenocarcinoma, NOS      0.57      0.40
0.47          10
                                         Adenocarcinoma, Papillary, NOS      0.00      0.00
0.00          1
                                         Non-Small Cell      1.00      1.00
1.00          1
                                         Squamous Cell Carcinoma, NOS      1.00      0.80
0.89          10

                                         accuracy
0.59          22
                                         macro avg      0.51      0.44
0.47          22
                                         weighted avg      0.76      0.59
0.66          22

Vote Model Results
Accuracy: 0.7727272727272727
Classification Report:

```

		precision	recall	f1-score
support				
10	Adenocarcinoma, NOS	0.80	0.80	0.80
1	Adenocarcinoma, Papillary, NOS	0.00	0.00	0.00
1	Non-Small Cell	0.00	0.00	0.00
10	Squamous Cell Carcinoma, NOS	0.75	0.90	0.82
22	accuracy			0.77
22	macro avg	0.39	0.43	0.40
22	weighted avg	0.70	0.77	0.74

Part E. Pipelines

The sklearn pipeline function is a tool in machine learning that simplifies the workflow by encapsulating all the steps involved in a single object. It offers advantages such as simplicity, reproducibility, efficiency, flexibility, and integration. The Pipeline is built using a list of (key, value) pairs, where the key is a string containing the name you want to give this step and value is an estimator object (the method to be executed).

Pipeline Task

- Read the [Study Case notebook for a pipeline functions](#) to understand the principle of the pipeline function
- Implement a pipeline which prepares and classifies data
- Use a `GridSearchCV` object with the `Pipeline` object and a parameter grid to optimize choose the best hyper parameters

```
from DataProcessor import DataProcessor

pipeline = DataProcessor(input_file_exp =
'GSE58661_series_matrix.txt.gz', input_file_meta =
"Lung3.metadata.xls",
                        cols_to_grab = ['title',
                                         'characteristics.tag.gender', 'characteristics.tag.histology',
                                         'characteristics.tag.tumor.size.maximumdiameter'],
                        target='characteristics.tag.histology')

pipeline_finished, encoders = pipeline.pipeline_run()
print(encoders)
```

```

print(pipeline_finished.head())
print(pipeline.look_at_encode(encoders, "title"))

N component to reach 50% explained of the variance: 14
Explained variance list: [0.13625197 0.0657839 0.0472555 0.04508168
0.03637887 0.02773954
0.02481434 0.0213042 0.02095335 0.0170285 0.01591754 0.01557377
0.01436692 0.01387649]
['title', 'gender', 'histology']
{'title': LabelEncoder(), 'gender': LabelEncoder()}
    title   gender           histology
tumor.size.maximumdiameter \
0      0       1  Squamous Cell Carcinoma, NOS
4.0
1      11      1  Adenocarcinoma, Papillary, NOS
1.3
2      22      1           Non-Small Cell
11.0
4      43      0  Squamous Cell Carcinoma, NOS
7.8
5      54      1  Adenocarcinoma, NOS
3.5

          PCA_1     PCA_2     PCA_3     PCA_4     PCA_5     PCA_6
PCA_7 \
0  15.075265 -5.472475  0.144114  2.997496  3.194623 -0.863421 -
8.363304
1  -6.797404  9.108340 -4.146173 10.091751 -3.858601  2.189105 -
0.724275
2   7.656460  5.641157  2.396656  0.507307 -0.120188  7.554593 -
9.346953
4   6.930503 -1.981333 -3.090020 -2.861382 -4.963344  0.083323
3.793841
5  -6.866965  0.155863 -4.951329 -4.215494  0.506148 -1.350626
0.538385

          PCA_8     PCA_9     PCA_10    PCA_11    PCA_12    PCA_13
PCA_14
0  2.396658  0.470345 -2.003722 -3.817031  0.566312 -2.601190 -
1.235243
1  2.863347 -2.221778 -1.284008  5.740897 -0.271739  2.137453 -
2.982435
2  4.866749 -4.272516  5.607741 -9.701937 -2.217542 11.334382
9.983083
4 -4.423526  4.856949  1.907469  1.411168 -3.258795  2.172437
4.477996
5 -2.566367 -0.201065  0.043919 -2.282812  2.594097  4.914067
0.438946

```

```

{'lung_1': 0, 'lung_10': 1, 'lung_11': 2, 'lung_12': 3, 'lung_13': 4,
'lung_14': 5, 'lung_15': 6, 'lung_16': 7, 'lung_17': 8, 'lung_18': 9,
'lung_19': 10, 'lung_2': 11, 'lung_20': 12, 'lung_21': 13, 'lung_22':
14, 'lung_23': 15, 'lung_24': 16, 'lung_25': 17, 'lung_26': 18,
'lung_27': 19, 'lung_28': 20, 'lung_29': 21, 'lung_3': 22, 'lung_30':
23, 'lung_31': 24, 'lung_32': 25, 'lung_33': 26, 'lung_34': 27,
'lung_35': 28, 'lung_36': 29, 'lung_37': 30, 'lung_38': 31, 'lung_39':
32, 'lung_40': 33, 'lung_41': 34, 'lung_42': 35, 'lung_43': 36,
'lung_44': 37, 'lung_45': 38, 'lung_46': 39, 'lung_47': 40, 'lung_48':
41, 'lung_49': 42, 'lung_5': 43, 'lung_50': 44, 'lung_51': 45,
'lung_52': 46, 'lung_53': 47, 'lung_54': 48, 'lung_55': 49, 'lung_56':
50, 'lung_57': 51, 'lung_58': 52, 'lung_59': 53, 'lung_6': 54,
'lung_60': 55, 'lung_61': 56, 'lung_62': 57, 'lung_63': 58, 'lung_64':
59, 'lung_65': 60, 'lung_66': 61, 'lung_67': 62, 'lung_68': 63,
'lung_69': 64, 'lung_7': 65, 'lung_70': 66, 'lung_71': 67, 'lung_72':
68, 'lung_73': 69, 'lung_74': 70, 'lung_75': 71, 'lung_76': 72,
'lung_77': 73, 'lung_78': 74, 'lung_79': 75, 'lung_8': 76, 'lung_80':
77, 'lung_81': 78, 'lung_82': 79, 'lung_83': 80, 'lung_84': 81,
'lung_85': 82, 'lung_86': 83, 'lung_87': 84, 'lung_88': 85, 'lung_89':
86, 'lung_9': 87}
{'lung_1': 0, 'lung_10': 1, 'lung_11': 2, 'lung_12': 3, 'lung_13': 4,
'lung_14': 5, 'lung_15': 6, 'lung_16': 7, 'lung_17': 8, 'lung_18': 9,
'lung_19': 10, 'lung_2': 11, 'lung_20': 12, 'lung_21': 13, 'lung_22':
14, 'lung_23': 15, 'lung_24': 16, 'lung_25': 17, 'lung_26': 18,
'lung_27': 19, 'lung_28': 20, 'lung_29': 21, 'lung_3': 22, 'lung_30':
23, 'lung_31': 24, 'lung_32': 25, 'lung_33': 26, 'lung_34': 27,
'lung_35': 28, 'lung_36': 29, 'lung_37': 30, 'lung_38': 31, 'lung_39':
32, 'lung_40': 33, 'lung_41': 34, 'lung_42': 35, 'lung_43': 36,
'lung_44': 37, 'lung_45': 38, 'lung_46': 39, 'lung_47': 40, 'lung_48':
41, 'lung_49': 42, 'lung_5': 43, 'lung_50': 44, 'lung_51': 45,
'lung_52': 46, 'lung_53': 47, 'lung_54': 48, 'lung_55': 49, 'lung_56':
50, 'lung_57': 51, 'lung_58': 52, 'lung_59': 53, 'lung_6': 54,
'lung_60': 55, 'lung_61': 56, 'lung_62': 57, 'lung_63': 58, 'lung_64':
59, 'lung_65': 60, 'lung_66': 61, 'lung_67': 62, 'lung_68': 63,
'lung_69': 64, 'lung_7': 65, 'lung_70': 66, 'lung_71': 67, 'lung_72':
68, 'lung_73': 69, 'lung_74': 70, 'lung_75': 71, 'lung_76': 72,
'lung_77': 73, 'lung_78': 74, 'lung_79': 75, 'lung_8': 76, 'lung_80':
77, 'lung_81': 78, 'lung_82': 79, 'lung_83': 80, 'lung_84': 81,
'lung_85': 82, 'lung_86': 83, 'lung_87': 84, 'lung_88': 85, 'lung_89':
86, 'lung_9': 87}

from Evaluation import Evaluator
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

svcl = SVC(kernel='sigmoid', probability=True)
knncl = KNeighborsClassifier(n_neighbors=5)
rfcl = RandomForestClassifier(random_state=29)

```

```

models = [svcl, knncl, rfcl]

y = pipeline_finished['histology']
X = pipeline_finished.drop(columns=['histology', 'title', 'gender'])

eval_class = Evaluator(models, X, y)

eval_class.execute()

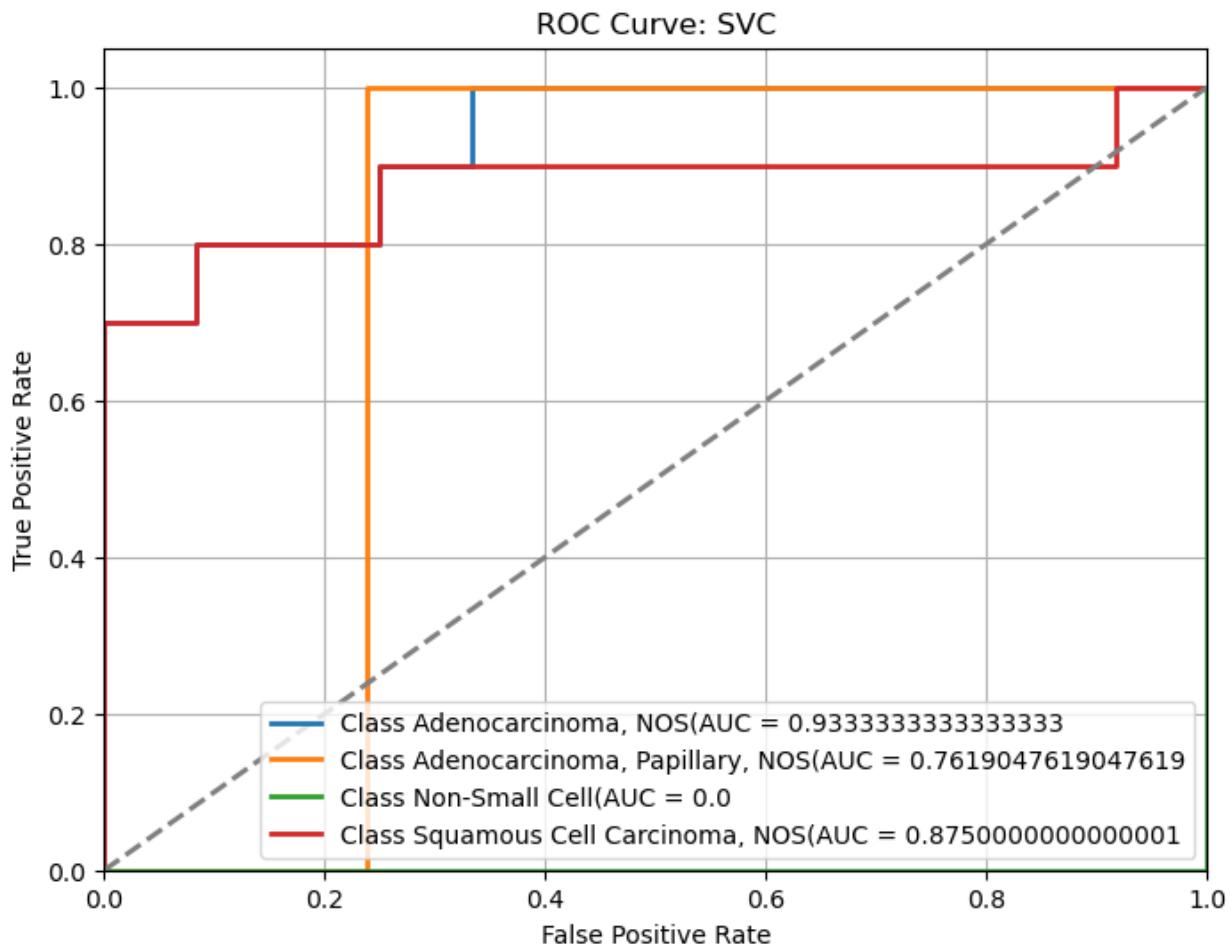
```

Lost X: 15/88 instances.
 Lost X: 15/88 instances.

Model name: SVC
 Precision: 0.7045454545454546
 Accuracy: 0.7727272727272727
 Recall: 0.7727272727272727
 F1 Score: 0.7355371900826445
 ConfusionMatrix:
 Matrix: [[9 0 0 1]
 [1 0 0 0]
 [0 0 0 1]
 [2 0 0 8]]

Classification report:

	Report	precision	recall	f1-score
	support			
10	Adenocarcinoma, NOS	0.75	0.90	0.82
1	Adenocarcinoma, Papillary, NOS	0.00	0.00	0.00
1	Non-Small Cell	0.00	0.00	0.00
10	Squamous Cell Carcinoma, NOS	0.80	0.80	0.80
22	accuracy			0.77
22	macro avg	0.39	0.43	0.40
22	weighted avg	0.70	0.77	0.74

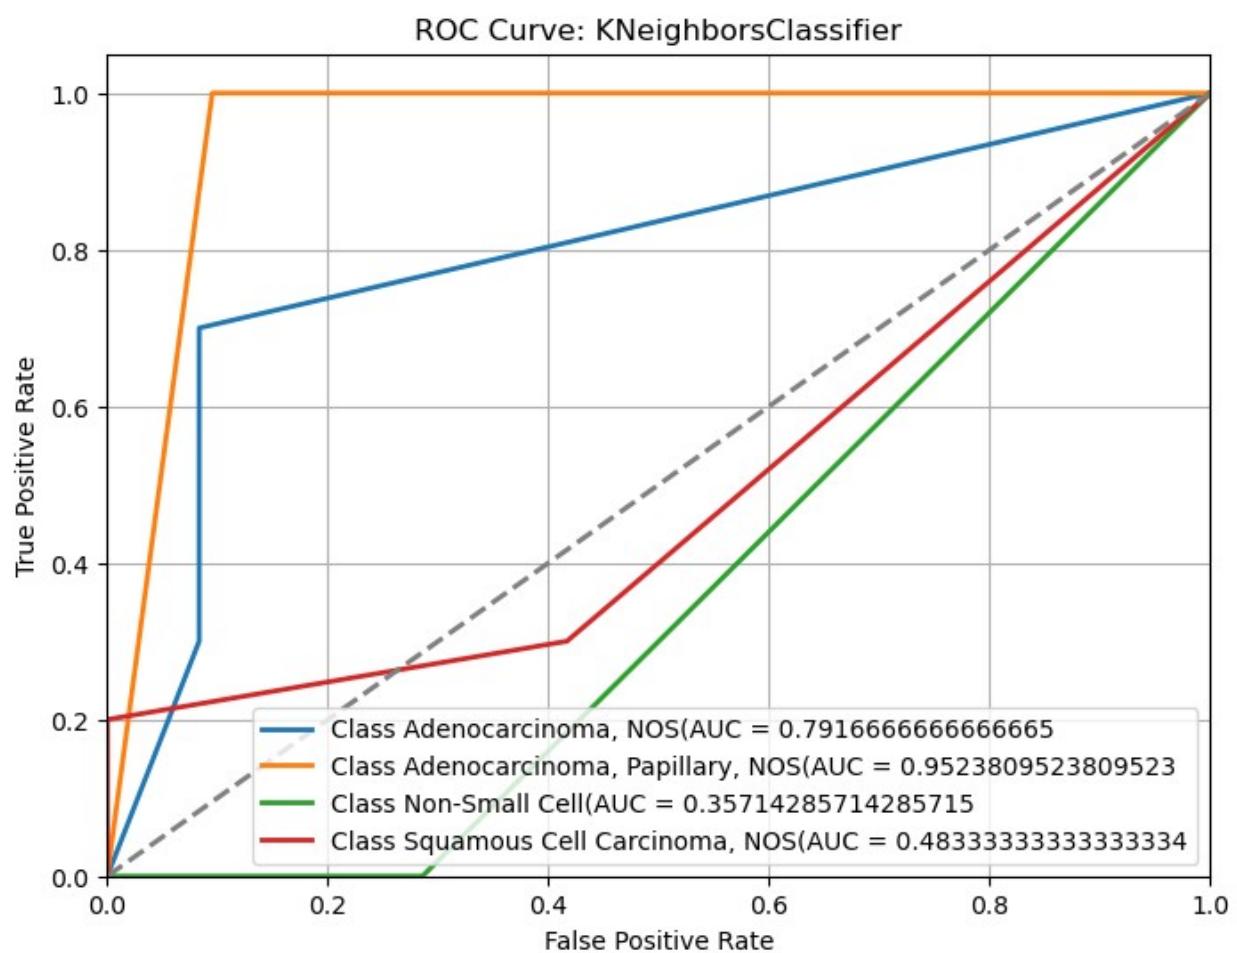


```

Model name: KNeighborsClassifier
Precision: 0.7196969696969697
Accuracy: 0.6363636363636364
Recall: 0.6363636363636364
F1 Score: 0.65599173553719
ConfusionMatrix:
Matrix: [[0 0 0 0 0]
 [3 5 0 0 2]
 [0 1 0 0 0]
 [0 0 0 0 1]
 [1 0 0 0 9]]
Classification report:
Report                                                 precision
recall   f1-score   support
Adenocarcinoma, Bronchiolo-alveolar Features        0.00      0.00
0.00          0
Adenocarcinoma, NOS                           0.83      0.50

```

0.62	10	Adenocarcinoma, Papillary, NOS	0.00	0.00
0.00	1	Non-Small Cell	0.00	0.00
0.00	1	Squamous Cell Carcinoma, NOS	0.75	0.90
0.82	10			
accuracy				
0.64	22		macro avg	0.32
0.29	22		weighted avg	0.72
0.66	22			0.64



Model name: RandomForestClassifier

Precision: 0.6942148760330579

Accuracy: 0.7272727272727273

Recall: 0.7272727272727273

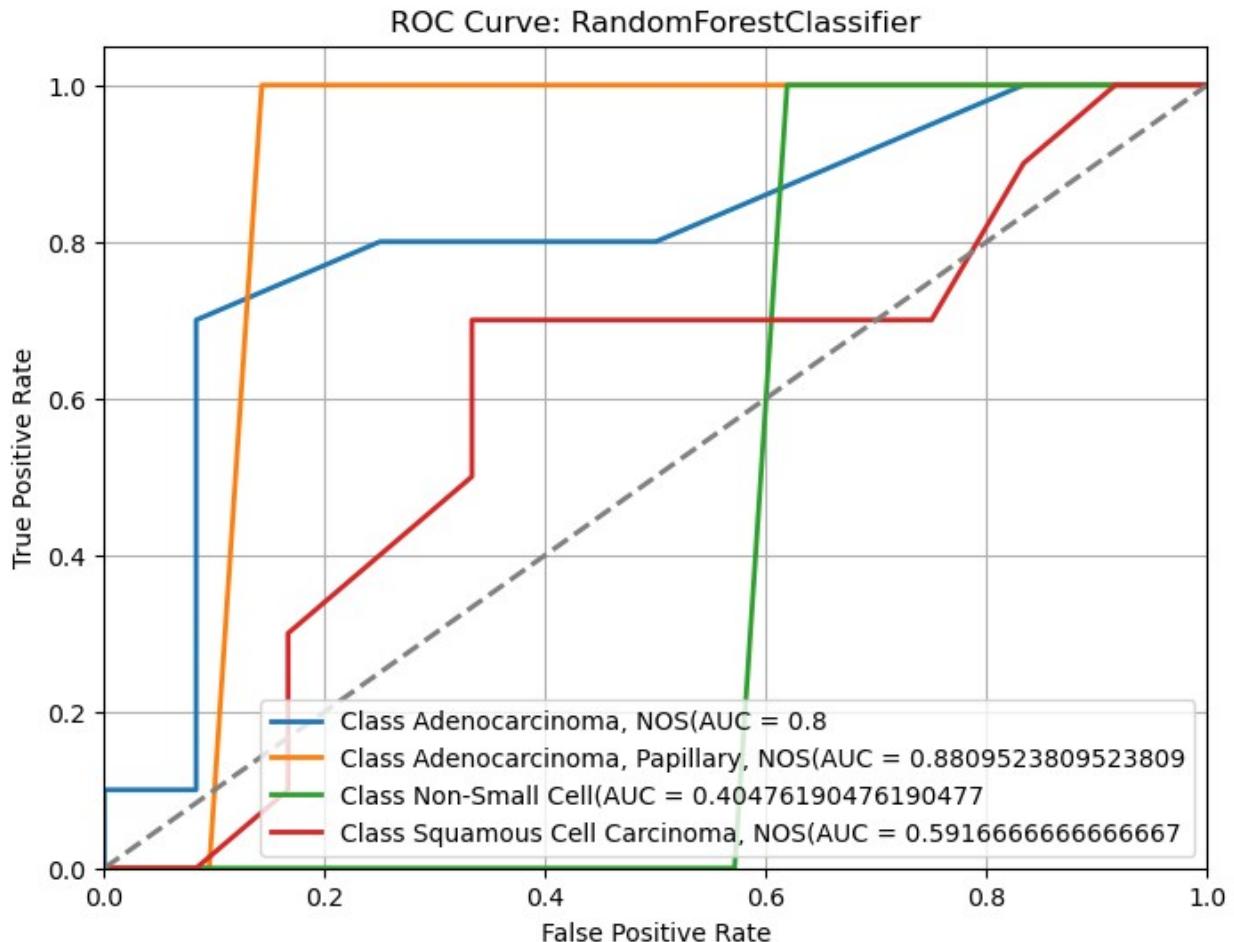
F1 Score: 0.7099567099567099

ConfusionMatrix:

Matrix: [[7 0 0 3]
[1 0 0 0]
[0 0 1 0]
[2 0 0 8]]

Classification report:

		precision	recall	f1-score
	support			
10	Adenocarcinoma, NOS	0.70	0.70	0.70
1	Adenocarcinoma, Papillary, NOS	0.00	0.00	0.00
1	Non-Small Cell	1.00	1.00	1.00
10	Squamous Cell Carcinoma, NOS	0.73	0.80	0.76
22	accuracy			0.73
22	macro avg	0.61	0.62	0.62
22	weighted avg	0.69	0.73	0.71



```

'Adenocarcinoma, Bronchiolo-alveolar Features',
'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
'Adenocarcinoma, Bronchiolo-alveolar Features',
'Squamous Cell Carcinoma, NOS',
'Adenocarcinoma, Bronchiolo-alveolar Features',
'Adenocarcinoma, Bronchiolo-alveolar Features'],
dtype=object),
<Figure size 800x600 with 1 Axes>],
'RandomForestClassifier': [RandomForestClassifier(random_state=29),
array(['Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
'Non-Small Cell', 'Adenocarcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Adenocarcinoma, NOS',
'Adenocarcinoma, NOS', 'Squamous Cell Carcinoma, NOS',
'Squamous Cell Carcinoma, NOS', 'Squamous Cell Carcinoma,
NOS',
'Adenocarcinoma, NOS', 'Adenocarcinoma, NOS'], dtype=object),
<Figure size 800x600 with 1 Axes>]}

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from Evaluation import Evaluator
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from VoteClass import Voter
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report

svccl = SVC(kernel='sigmoid', probability=True)
knncl = KNeighborsClassifier(n_neighbors=5)
rfcl = RandomForestClassifier(random_state=29)

models = [svccl, knncl, rfcl]

```

```

y = pipeline_finished['histology'].copy()
X = pipeline_finished.drop(columns=['histology', 'title', 'gender'])

rare_classes = y.value_counts()[y.value_counts() < 4].index # less
than 3 samples
y = y.replace(rare_classes, 'Other')

# Find minimum samples in a class
min_class_samples = y.value_counts().min()
print("Minimum samples in a class:", min_class_samples)

# Use this for cv
cv_strategy = StratifiedKFold(n_splits=min(min_class_samples, 5))

eval_class = Evaluator(models, X, y)
eval_class.get_train_test_split()

def pipeline_acceptor(model_pipeline, model_param_grid, cv_strategy,
eval_class):
    model_grid = GridSearchCV(model_pipeline, model_param_grid,
cv=cv_strategy, scoring='accuracy')
    model_grid.fit(eval_class.X_train, eval_class.y_train)
    model_name = model_pipeline.steps[-1][0].upper()
    print(f"Best {model_name} params:", model_grid.best_params_)
    print(f"Best {model_name} score:", model_grid.best_score_)
    return model_grid.best_estimator_

# Pipeline + GridSearchCV for SVC
svc_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('svc', SVC(probability=True))
])

svc_param_grid = {
    'svc_kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
    'svc_C': [0.1, 1, 10, 100],
    'svc_gamma': ['scale', 'auto']
}

# Pipeline + GridSearchCV for KNN
knn_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('knn', KNeighborsClassifier())
])

```

```

knn_param_grid = {
    'knn_n_neighbors': [3, 5, 7, 9],
    'knn_weights': ['uniform', 'distance'],
    'knn_p': [1, 2] # 1=Manhattan, 2=Euclidean
}

# Pipeline + GridSearchCV for Random Forest
rf_pipeline = Pipeline([
    ('rf', RandomForestClassifier(random_state=29))
])

rf_param_grid = {
    'rf_n_estimators': [50, 100, 200],
    'rf_max_depth': [None, 5, 10, 20],
    'rf_min_samples_split': [2, 5, 10],
    'rf_min_samples_leaf': [1, 2, 4]
}

model_list2 = [pipeline_acceptor(rf_pipeline, rf_param_grid,
cv_strategy, eval_class), pipeline_acceptor(knn_pipeline,
knn_param_grid, cv_strategy, eval_class),
pipeline_acceptor(svc_pipeline, svc_param_grid, cv_strategy,
eval_class)]

vote = Voter(model_list=model_list2, X_test=eval_class.X_test,
vote_type='s')
results = vote.choose_vote()

print("Vote Model Results")
print("Accuracy:", accuracy_score(eval_class.y_test, results))
print("Classification Report:\n",
classification_report(eval_class.y_test, results, zero_division=0))

Minimum samples in a class: 4

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn()

Best RF params: {'rf_max_depth': 5, 'rf_min_samples_leaf': 1,
'rf_min_samples_split': 5, 'rf_n_estimators': 50}
Best RF score: 0.6229166666666667

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn()

```

```

Best KNN params: {'knn_n_neighbors': 5, 'knn_p': 1, 'knn_weights': 'distance'}
Best KNN score: 0.6229166666666668

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn()

Best SVC params: {'svc_C': 10, 'svc_gamma': 'scale', 'svc_kernel': 'rbf'}
Best SVC score: 0.6541666666666667
Vote Model Results
Accuracy: 0.37037037037037035
Classification Report:
               precision    recall
f1-score   support
Adenocarcinoma, Bronchiolo-alveolar Features      0.00     0.00
0.00          3
                           Adenocarcinoma, NOS      0.40     0.25
0.31          8
                           Adenocarcinoma, Papillary, NOS  0.00     0.00
0.00          3
                           Non-Small Cell      0.00     0.00
0.00          3
                           Other          0.20     0.33
0.25          3
                           Squamous Cell Carcinoma, NOS  0.41     1.00
0.58          7

               accuracy
0.37          27
                           macro avg      0.17     0.26
0.19          27
                           weighted avg  0.25     0.37
0.27          27

from PipelineOptimize import PipelineMaker

y = pipeline_finished['histology'].copy()
X = pipeline_finished.drop(columns=['histology', 'title', 'gender'])

# Pipeline + GridSearchCV for SVC
svc_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('svc', SVC(probability=True))
])

svc_param_grid = {

```

```

'svc_kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
'svc_C': [0.1, 1, 10, 100],
'svc_gamma': ['scale', 'auto']
}

# Pipeline + GridSearchCV for KNN
knn_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('knn', KNeighborsClassifier())
])

knn_param_grid = {
    'knn_n_neighbors': [3, 5, 7, 9],
    'knn_weights': ['uniform', 'distance'],
    'knn_p': [1, 2] # 1=Manhattan, 2=Euclidean
}

# Pipeline + GridSearchCV for Random Forest
rf_pipeline = Pipeline([
    ('rf', RandomForestClassifier(random_state=29))
])

rf_param_grid = {
    'rf_n_estimators': [50, 100, 200],
    'rf_max_depth': [None, 5, 10, 20],
    'rf_min_samples_split': [2, 5, 10],
    'rf_min_samples_leaf': [1, 2, 4]
}

pipeline_grid_list = [[svc_pipeline, svc_param_grid], [rf_pipeline,
rf_param_grid], [knn_pipeline, knn_param_grid]]

piper = PipelineMaker(X=X,y=y, pipeline_grid_list=pipeline_grid_list)
best_model_list = piper.execute_all()

print(best_model_list)

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn

Best SVC params: {'svc_C': 10, 'svc_gamma': 'scale', 'svc_kernel':
'rbf'}
Best SVC score: 0.6541666666666667

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn

```

```

Best RF params: {'rf__max_depth': 5, 'rf__min_samples_leaf': 1,
'rf__min_samples_split': 5, 'rf__n_estimators': 50}
Best RF score: 0.6229166666666667

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn()

Best KNN params: {'knn__n_neighbors': 5, 'knn__p': 1, 'knn__weights':
'distance'}
Best KNN score: 0.6229166666666668
[Pipeline(steps=[('scaler', StandardScaler()),
                 ('svc', SVC(C=10, probability=True))]),
 Pipeline(steps=[('rf',
                  RandomForestClassifier(max_depth=5,
                                         min_samples_split=5,
                                         n_estimators=50,
                                         random_state=29))], Pipeline(steps=[('scaler', StandardScaler()),
                                              ('knn', KNeighborsClassifier(p=1,
                                         weights='distance'))])]

piper_solo = PipelineMaker(X=X, y=y, model_pipeline=svc_pipeline,
model_param_grid=svc_param_grid)
best_model = piper_solo.execute_all()
print(best_model)

/home/jippe/anaconda3/envs/dsc5/lib/python3.13/site-packages/sklearn/
model_selection/_split.py:811: UserWarning: The least populated class
in y has only 1 members, which is less than n_splits=4.
warnings.warn()

Best SVC params: {'svc__C': 10, 'svc__gamma': 'scale', 'svc__kernel':
'rbf'}
Best SVC score: 0.6541666666666667
Pipeline(steps=[('scaler', StandardScaler()),
                 ('svc', SVC(C=10, probability=True))])

```

Bring it all together

By now, you've developed code snippets for model evaluation, optimization, and data improvement. Now, leverage these skills to build a classification model using the Clinical and Genetic Lung data. Make sure that you log your experiments.

Once you're satisfied with the model, upload the relevant code to your repository and or refactor code with new insights. Furthermore, take a moment to reflect on its applicability and potential real-world impact. Update your evaluation document(s) with these findings in your repository.

Bonus: ML for operations

If we intend to deploy the model in a real-world application, it's more efficient to save and reuse the trained model rather than retraining it each time.

- Read the blog: <https://neptune.ai/blog/saving-trained-model-in-python>
- Write three python files 1) a train_model python file 2) a use_model python file 3) a retrain_model python file that adds new data to the original training data and updates the model
- Update your repository